

Tutorial on amortized optimization for learning to optimize over continuous spaces

Brandon Amos, Meta AI (FAIR)



[brandondamos](https://twitter.com/brandondamos)



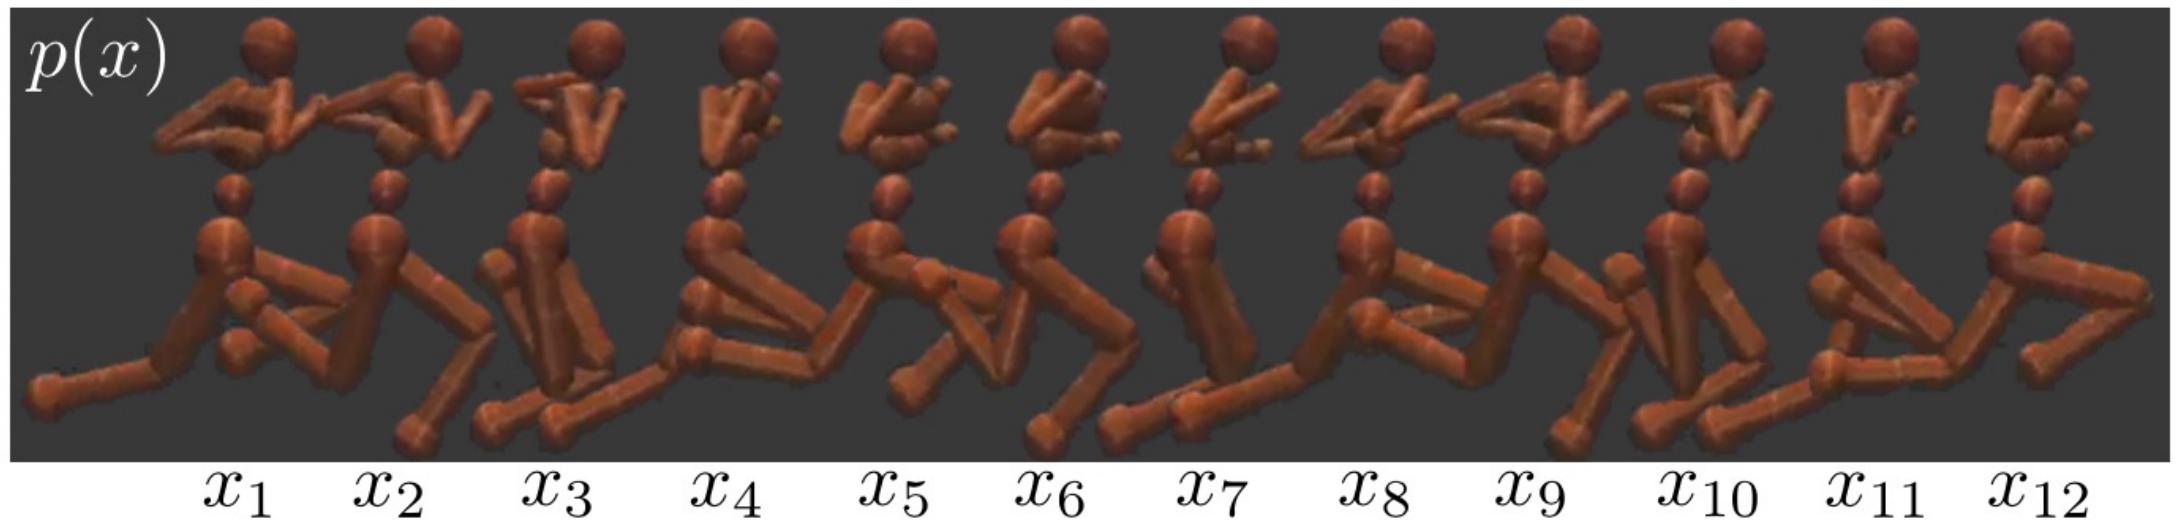
[bamos.github.io](https://github.com/bamos)

Optimization is a powerful modeling tool

Continuous optimization expresses many **non-trivial operations**

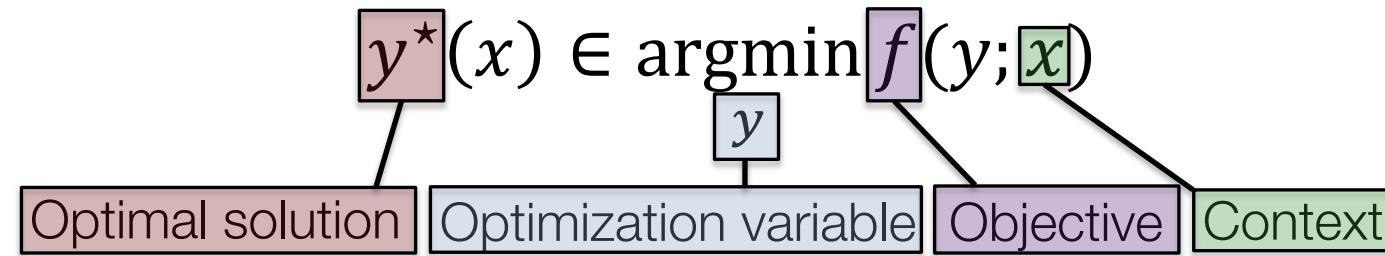
Control, reinforcement learning, robotics, geometry (projections), variational inference, finance (portfolio optimization), sparse coding, meta-learning, deep equilibrium networks, optimal transport, game and market equilibrium

Application: Continuous control of a robotic system:

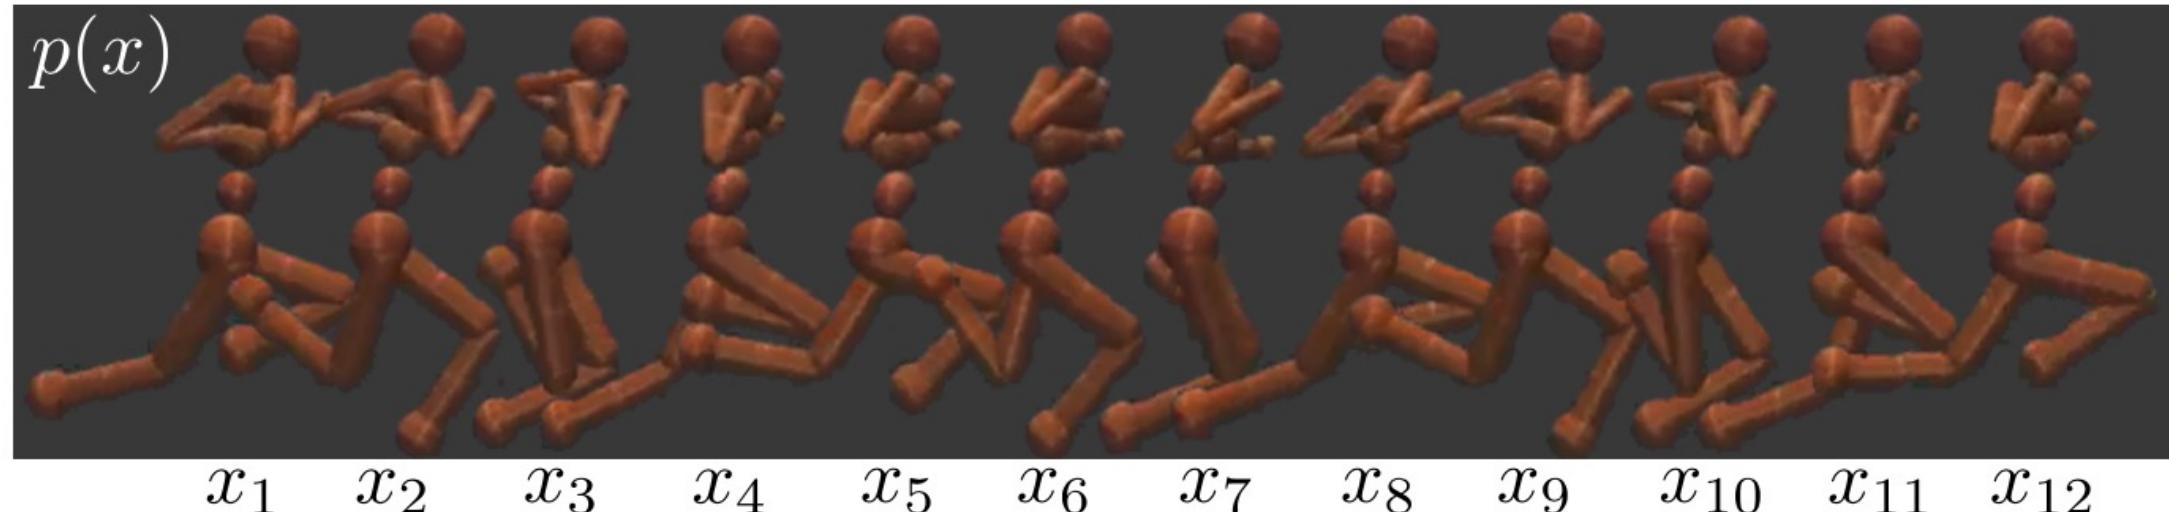


Repeatedly solving optimization problems

Optimization problems often **do not live in isolation** and are often **repeatedly solved** in deployment



In control, x is the system state, y is the action, $f(y; x)$ is a cost, and $y^*(x)$ is an optimal action



Difficulty: Optimization is computationally expensive

Sometimes solving just **once** may be difficult

Exasperated when **repeatedly solving** during deployment

Insight: Shared structure between problems

Problems **share structure** and don't live in isolation

Solution: Amortized optimization

Use **machine learning** to uncover the shared structure

Create **learning-augmented** versions of classical optimization solvers

Far surpasses average or worst-case convergence rates

Amortized optimization foundations

This talk: Explore foundations and connect applications

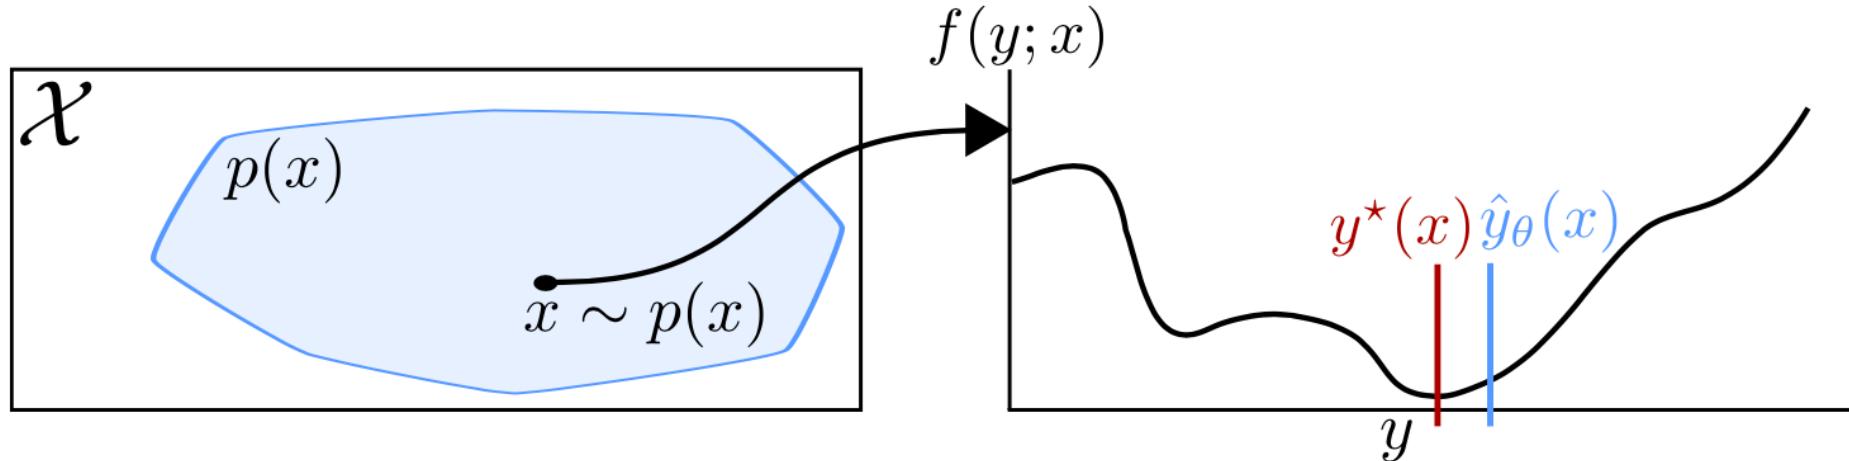
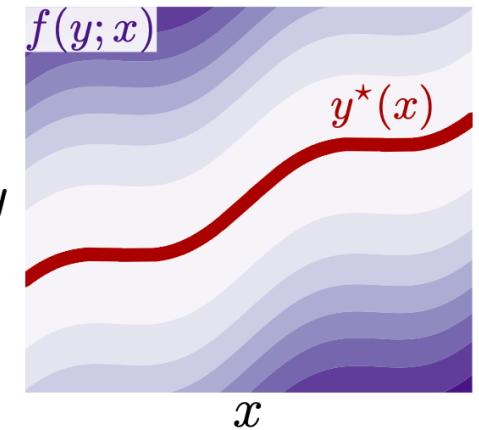
Amortization model $\hat{y}_\theta(x)$ tries to approximate $y^*(x)$

Example: A neural network mapping from x to the solution

Loss \mathcal{L} measures how well \hat{y} fits y^* and optimized with $\min_{\theta} \mathcal{L}(\hat{y}_\theta)$

Example: $\mathcal{L}(\hat{y}_\theta) := \mathbb{E}_{p(x)} \|\hat{y}_\theta(x) - y^*(x)\|_2^2$

$$y^*(x) \in \operatorname{argmin}_y f(y; x)$$



This talk: Amortized optimization

Motivation and context

Design decisions

Modeling paradigms for \hat{y}_θ (fully-amortized and semi-amortized models)

Learning paradigms for \mathcal{L} (objective-based and regression-based)

Applications

VAEs, hyper-networks, MAML, reinforcement learning, meta optimal transport

Extensions and limitations

This talk: Amortized optimization

Motivation and context

Design decisions

Modeling paradigms for \hat{y}_θ (fully-amortized and semi-amortized models)

Learning paradigms for \mathcal{L} (objective-based and regression-based)

Applications

VAEs, hyper-networks, MAML, reinforcement learning, meta optimal transport

Extensions and limitations

Modeling paradigms for \hat{y}_θ

Fully-amortized models: Map from the context x to the solution **without** accessing the objective f

Example: Neural network mapping from x to the solution

Most of our applications will focus on these

Semi-amortized models: Internally access the objective f

Example: Gradient-based meta-learning models such as MAML

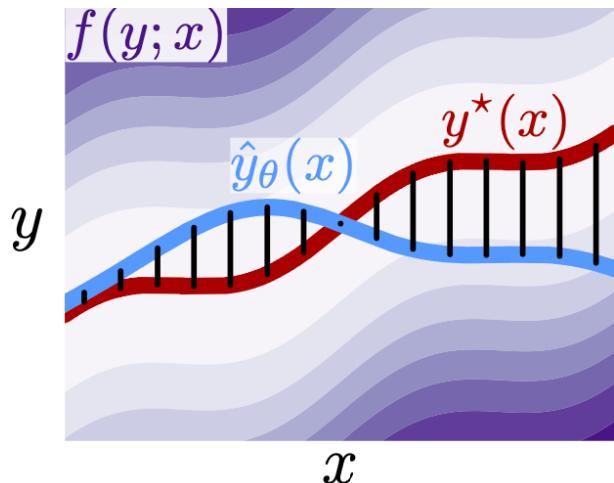
$$\hat{y}_\theta^0 \rightarrow \hat{y}_\theta^1 \rightarrow \dots \rightarrow \hat{y}_\theta^K =: \hat{y}_\theta(x)$$

Learning paradigms for \mathcal{L}

Regression-based

$$\mathcal{L}_{\text{reg}}(\hat{y}_\theta) := \mathbb{E}_{p(x)} \|\hat{y}_\theta(x) - y^*(x)\|_2^2$$

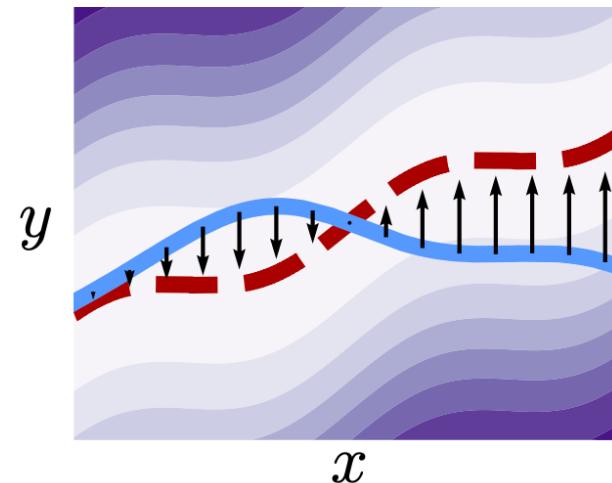
- Does not consider $f(y; x)$
- + Uses global information with $y^*(x)$
- Expensive to compute $y^*(x)$
- + Does not compute $\nabla_y f(y; x)$
- Hard to learn non-unique $y^*(x)$



Objective-based:

$$\mathcal{L}_{\text{obj}}(\hat{y}_\theta) := \mathbb{E}_{p(x)} f(\hat{y}_\theta(x); x)$$

- + Uses objective information of $f(y; x)$
- Can get stuck in local optima of $f(y; x)$
- + Faster, does not require $y^*(x)$
- Often requires computing $\nabla_y f(y; x)$
- + Easily learns non-unique $y^*(x)$



This talk: Amortized optimization

Motivation and context

Design decisions

Modeling paradigms for \hat{y}_θ (fully-amortized and semi-amortized models)

Learning paradigms for \mathcal{L} (objective-based and regression-based)

Applications

VAEs, hyper-networks, MAML, reinforcement learning, meta optimal transport

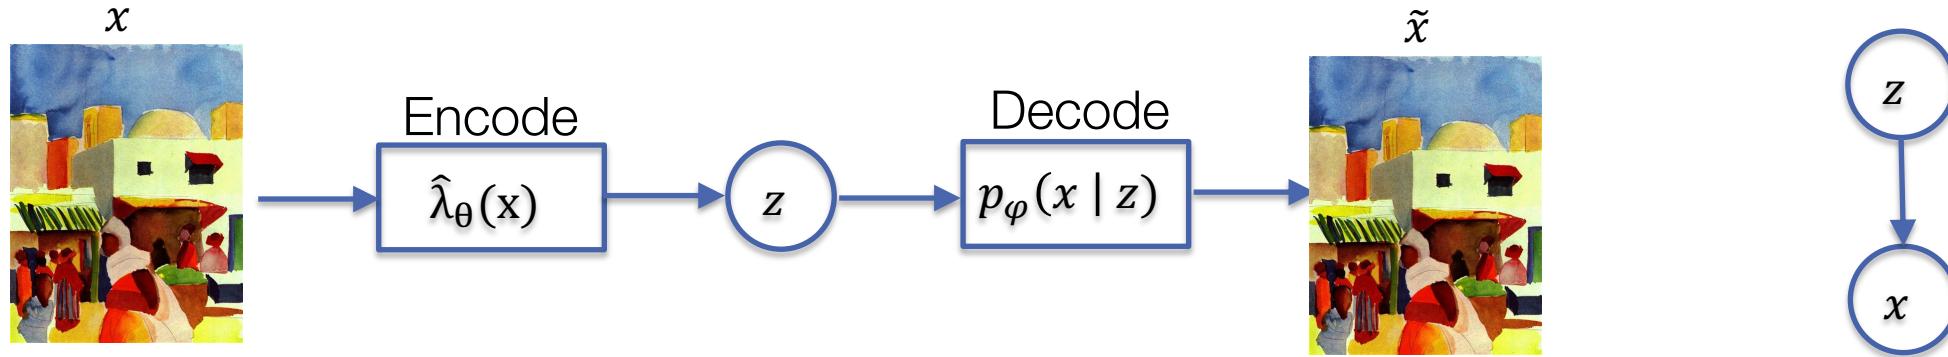
Extensions and limitations

Warning: Quick overview of separate applications

Goal: Give a broad overview of **existing uses of amortized optimization**

Will try to focus on the amortization components, but assumes some context

Amortized variational inference and auto-encoders



VAEs amortize the evidence lower bound (ELBO)

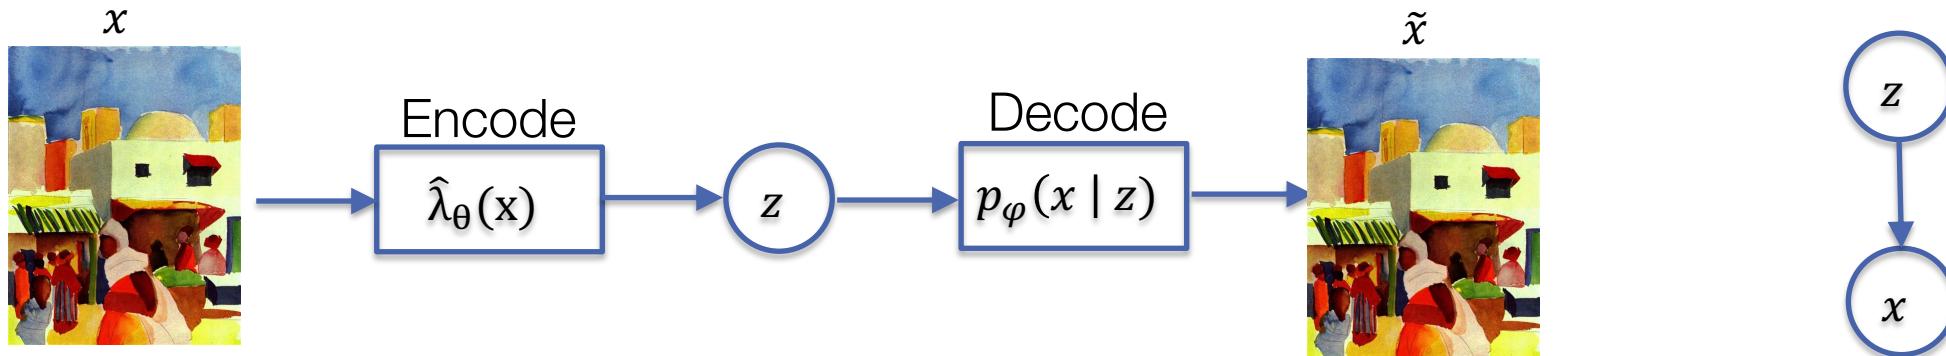
model: fully-amortized $\hat{\lambda}_\theta(x)$ as a neural network

loss: objective-based

objective: ELBO

learning: $\max_{\theta, \varphi} \mathbb{E}_{p(x)} \text{ELBO}(\hat{\lambda}_\theta(x); x, \varphi)$

Semi-amortized variational autoencoders



Incorporates knowledge of the ELBO into the encoder

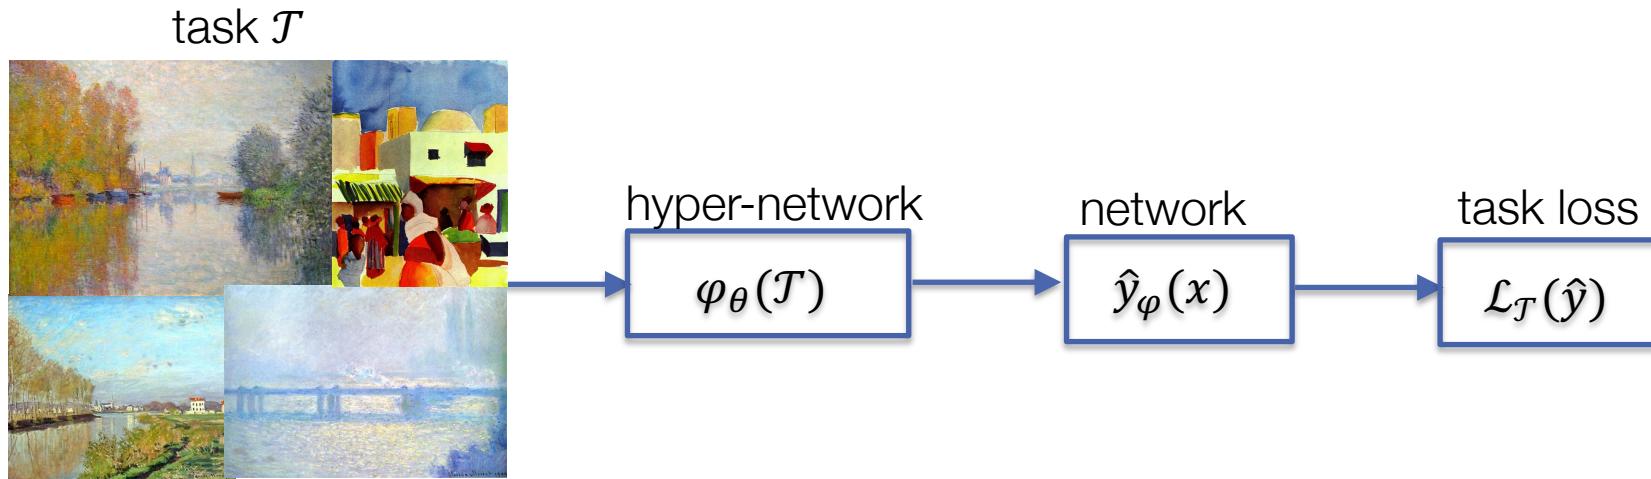
model: semi-amortized $\hat{\lambda}_\theta(x) \approx \underset{\lambda}{\operatorname{argmin}} \text{ELBO}(\lambda; x, \varphi)$, approximated by gradient updates

loss: objective-based

objective: ELBO (same as before)

learning: $\max_{\theta, \varphi} \mathbb{E}_{p(x)} \text{ELBO}(\hat{\lambda}_\theta(x); x, \varphi)$

Hyper-networks as amortized optimization



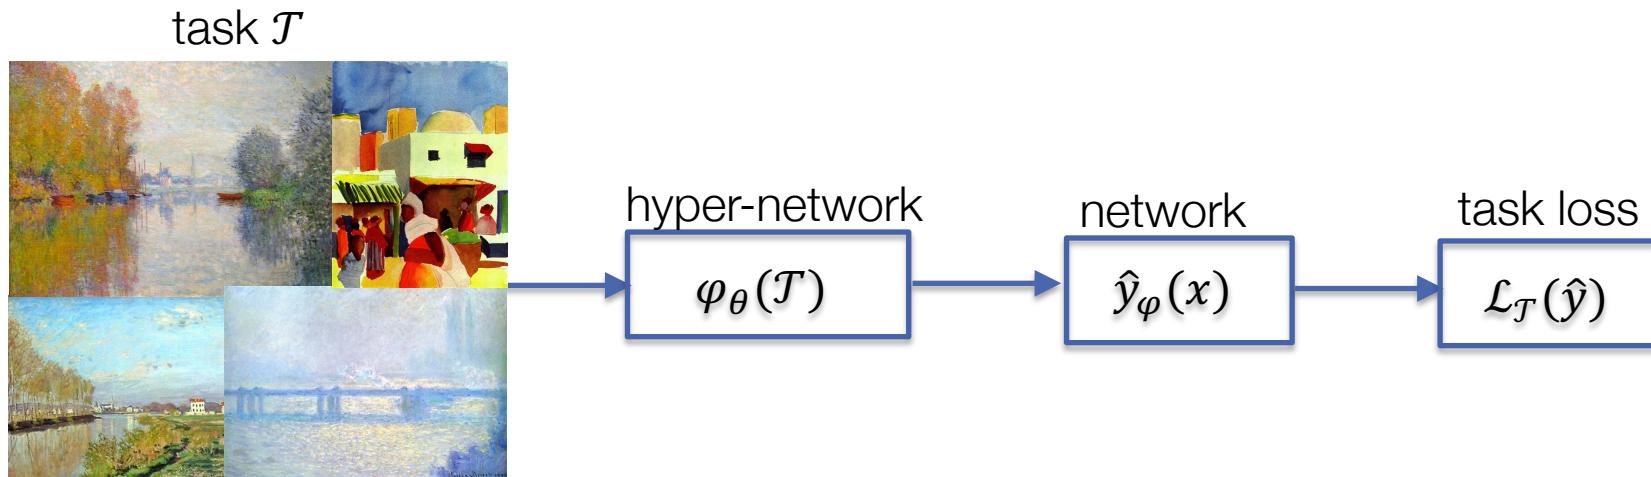
model: fully-amortized

loss: objective-based

objective: task loss

learning: $\operatorname{argmin}_{\theta} \mathbb{E}_{p(\mathcal{T})} \mathcal{L}_{\mathcal{T}}(\hat{y})$

MAML as amortized optimization



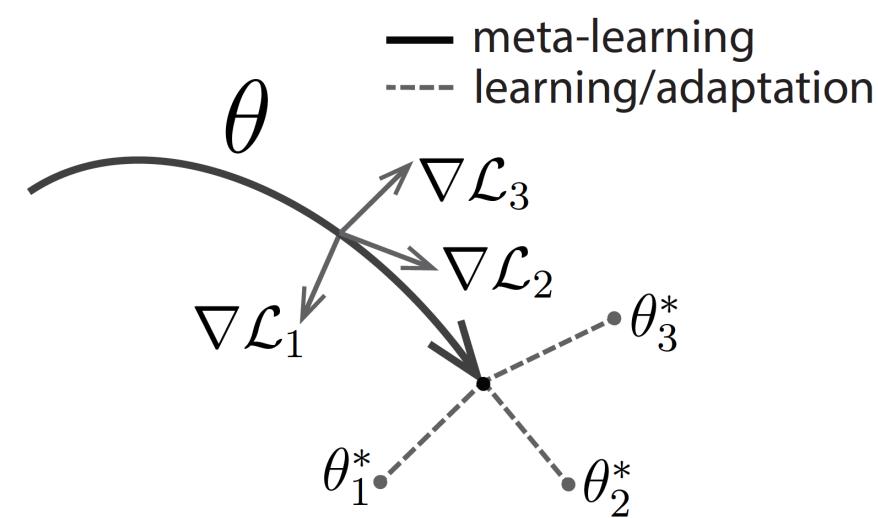
Setup almost identical to hyper-networks, except for the model

model: semi-amortized

loss: objective-based

objective: task loss

learning: $\min_{\theta} \mathbb{E}_{p(\mathcal{T})} \mathcal{L}_{\mathcal{T}}(\hat{y}_{\theta})$



Amortized policy learning in reinforcement learning

The **policy** $\pi_\theta(x)$ of almost every reinforcement learning method performs amortized optimization

setup: Controlling a continuous MDP

policy: $\pi_\theta(x)$

value (expected reward): $V(x) := \mathbb{E} r(x_t, u_t)$

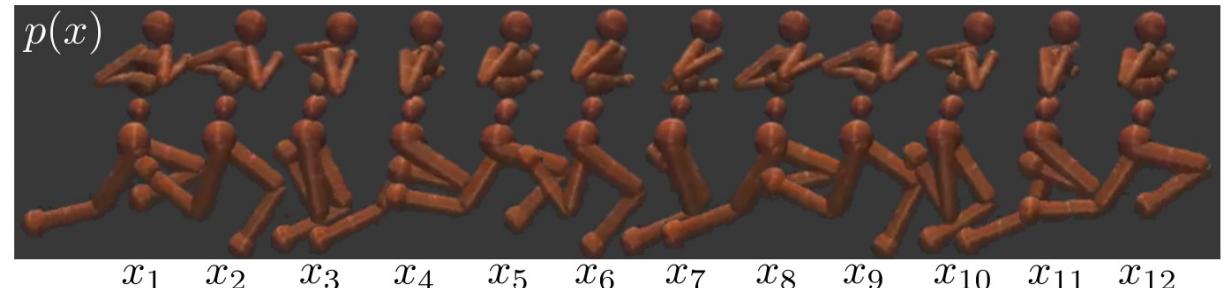
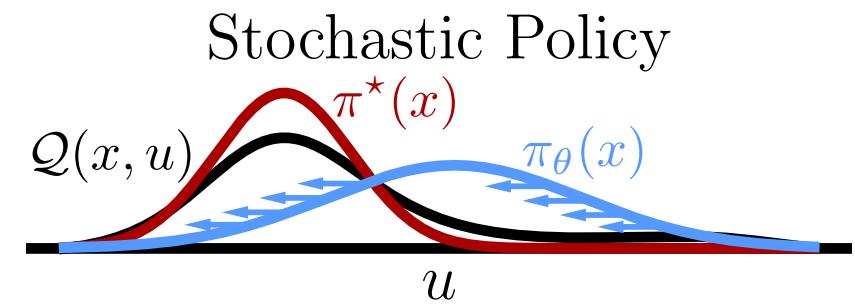
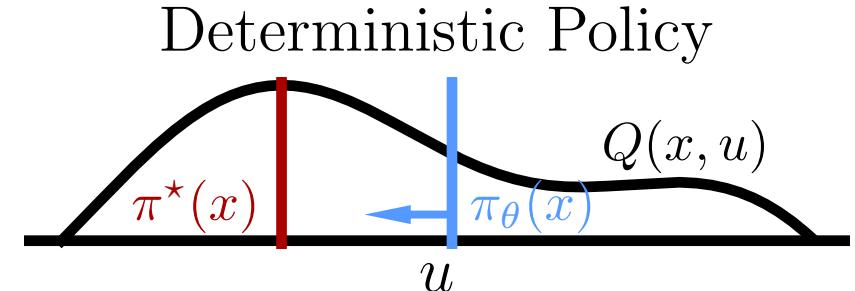
Q -value: $Q(x, u) := \sum r(x, u) + \mathbb{E}_{x'} V^\pi(x')$

model: fully-amortized

loss: objective-based

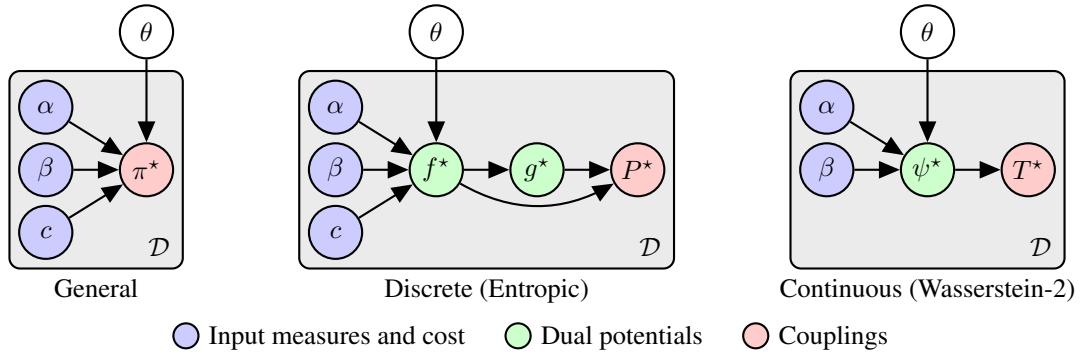
objective: Q -value

learning: $\max_{\theta} \mathbb{E}_{p(x)} Q(x, \pi_\theta(x))$



(Amortized) meta optimal transport

(Available on workplace)



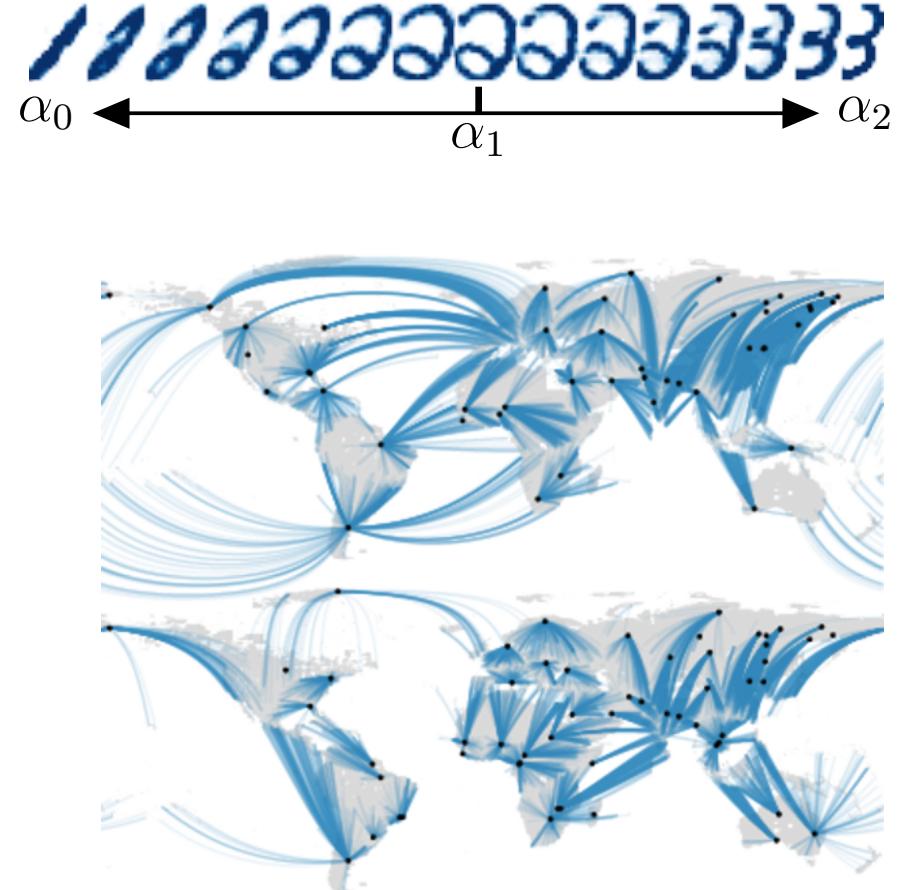
goal: Optimally transport mass between measures α to β

model: fully-amortized (hyper-network)

loss: objective-based

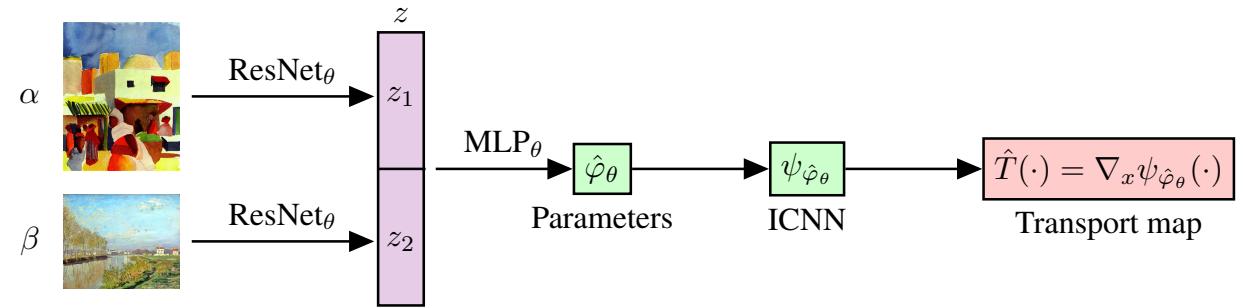
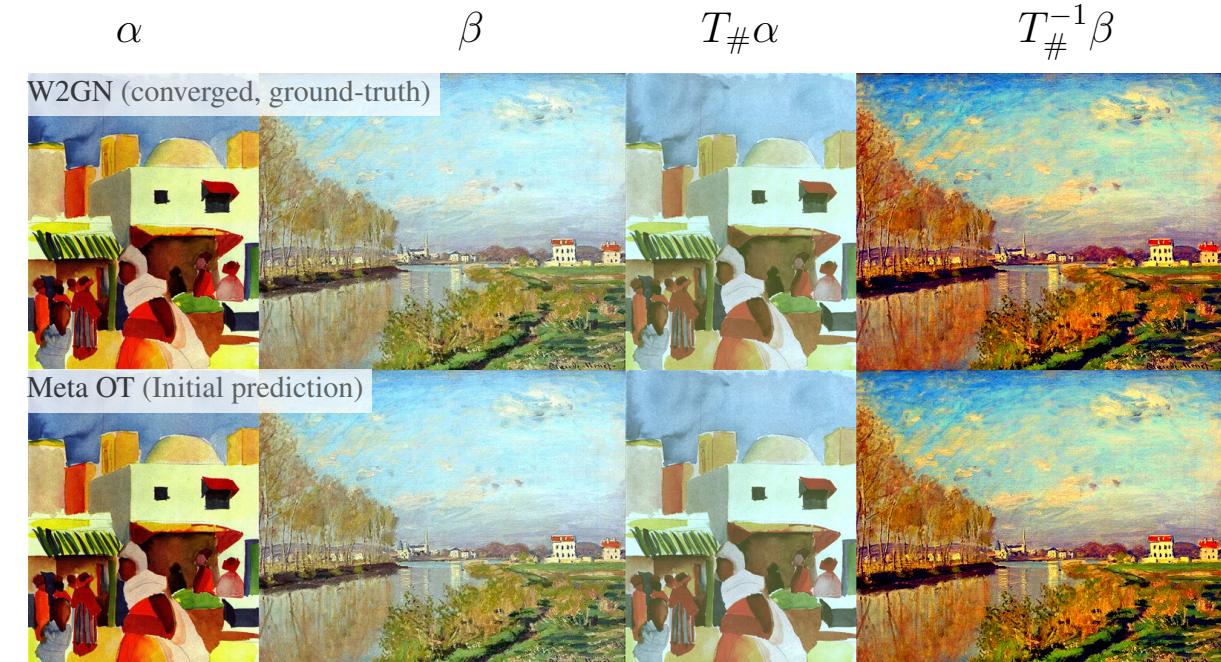
objective: transport cost

learning: $\max_{\theta} \mathbb{E}_{p(\alpha, \beta)} C(\hat{\pi}(\alpha, \beta))$



(Amortized) meta optimal transport

Predict couplings **10000x faster** than SOTA solvers that doesn't use learning



More meta color transfer predictions

α

β

$T_{\#}\alpha$

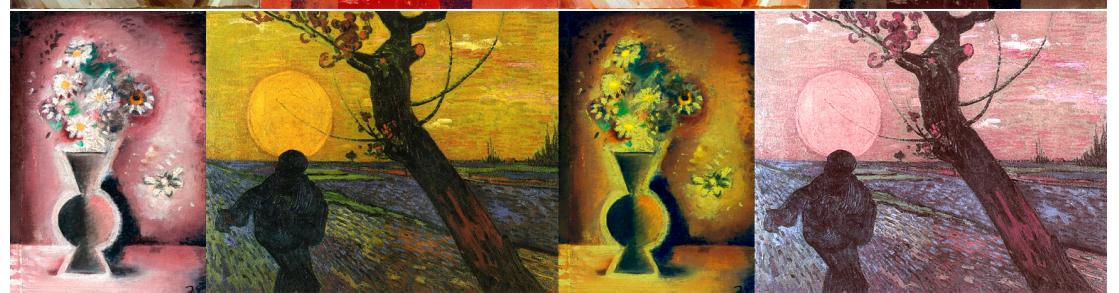
$T_{\#}^{-1}\beta$

α

β

$T_{\#}\alpha$

$T_{\#}^{-1}\beta$



This talk: Amortized optimization

Motivation and context

Design decisions

Modeling paradigms for \hat{y}_θ (fully-amortized and semi-amortized models)

Learning paradigms for \mathcal{L} (objective-based and regression-based)

Applications

VAEs, hyper-networks, MAML, reinforcement learning, meta optimal transport

Extensions and limitations

Limitations of amortized optimization

Gives **extreme computational advantages** in many settings
but may be **detrimental and suboptimal** in the wrong settings

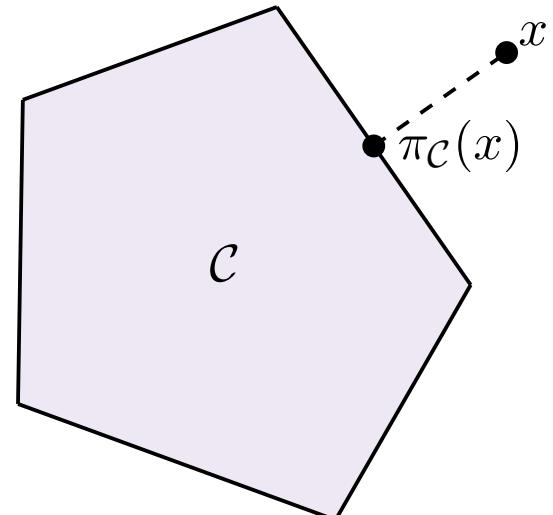
1. Difficult in **out-of-domain settings** when the contexts significantly change
2. Generally difficult to **ensure stability or convergence**
3. Typically **does not solve previously intractable problems**
4. Can be **difficult to obtain high-accuracy solutions** without fine-tuning/semi-amortization

Extensions

All settings here: **unconstrained continuous Euclidean optimization problems**

Possible to expand far beyond these:

1. **New applications and settings for semi-amortized modeling**
2. **Constrained domains** (e.g. with differentiable projections)
3. **Discrete optimization settings** (e.g. with differentiable discrete optimization)
4. **Non-Euclidean settings** (e.g. with Riemannian optimization)



Tutorial on amortized optimization for learning to optimize over continuous spaces

Brandon Amos, Meta AI (FAIR)



[brandondamos](https://twitter.com/brandondamos)



[bamos.github.io](https://github.com/bamos)

Motivation and context

Design decisions

Modeling paradigms for \hat{y}_θ (fully-amortized and semi-amortized models)

Learning paradigms for \mathcal{L} (objective-based and regression-based)

Applications

VAEs, hyper-networks, MAML, reinforcement learning, meta optimal transport

Extensions and limitations



github.com/facebookresearch/amortized-optimization-tutorial